

David Kieu

david.kieu25@gmail.com | linkedin.com/in/david-kieu-tech | github.com/monsieurkd | Adelaide, SA (open to relocation across AU)

Professional Summary

Data engineer who turns messy real-world data into reliable, tested pipelines and queryable stores. Built and validated automated data and ML pipelines (Python, SQL, PostgreSQL) and a computer-vision data pipeline for document images in a commercial setting, with strong SQL, data-validation, and CI/Docker discipline. CS student at the University of Adelaide (AI major, GPA 6.375/7.0), AWS Certified Cloud Practitioner, based in Adelaide and open to relocation across Australia.

Technical Skills

Languages	Python, SQL, Bash, JavaScript	Automation & Tooling	Reproducible Pipelines, NumPy
Databases	PostgreSQL, SQL Server, ChromaDB, Redis		Python Scripting, REST APIs, Task Scheduling
Data Engineering	ETL & Data Pipelines, Data Validation, Data Partitioning	DevOps & Cloud	Docker, AWS, CI/CD, Git, Linux

Education

University of Adelaide <i>Bachelor of Computer Science – Artificial Intelligence</i> <i>GPA: 6.375/7.0</i>	Adelaide, South Australia <i>2023 – 2026</i>
AWS Certified Cloud Practitioner <i>Amazon Web Services – Credential ID: AWS05628495</i>	December 2025

Work Experience

University of Adelaide <i>Research Intern – Machine Learning</i>	November 2024 – Present <i>Adelaide, Australia</i>
– Built a reproducible Python experimentation pipeline (PyTorch) that ran, logged, and validated results across 50+ model configurations, replacing manual, ad-hoc setup with scripted, re-runnable data workflows – Designed an automated training-and-evaluation pipeline with a curriculum-learning strategy that cut convergence time by 40%, producing systematic, validated outputs documented for research publication	
Kaopiz Software JSC <i>Machine Learning Engineer</i>	November 2023 – February 2024
– Built a Python and OpenCV image-processing pipeline to localise Japanese characters across a document-image dataset, scripting reproducible preprocessing from raw images to extracted character regions – Developed reusable data-preparation scripts for the team’s document-recognition R&D, replacing manual, ad-hoc image handling with repeatable processing	
Ecosmartvietnam <i>Software Engineer</i>	September 2022 – January 2023
– Architected and shipped a commercial full-stack application (Flask, PostgreSQL) with a normalised relational data model, owning the full SDLC end to end	

Projects

Autonomous Multi-Agent Data-Engineering Swarm Python	Multi-agent orchestration, Claude Code subagents,
– Built an autonomous multi-agent system that drives a data-engineering task end-to-end through role-specialised subagents (Planner, Builder, Verifier) as a plan-build-verify loop that returns a machine-checkable pass/fail. – Reliability comes from agent design, not prompting – each role has restricted tools (the Planner cannot edit code, the Verifier cannot write at all), and a file-based data contract is the sole hand-off between stages, so correctness is verifiable rather than left to the model.	